

Evolution of Strategies for Prisoner's Dilemma

Gleb Fedorovich, Jan Heldmann, Lukas
Schmitt, Fengshi Zheng





Historical background

Bob



Today



Tomorrow



Alice



Historical background



Bob

Today



Tomorrow



Alice



B Cooperates

B Defects

A Cooperates

A Defects

	A Cooperates	A Defects
B Cooperates	(3,3)	(0,5)
B Defects	(5,0)	(1,1)

Historical background



Bob



Today

Tomorrow



Alice

B Cooperates

B Defects

A Cooperates

A Defects

	A Cooperates	A Defects
B Cooperates	(3,3)	(0,5)
B Defects	(5,0)	(1,1)

Historical background



Bob



Today

Tomorrow



Alice



B Cooperates

B Defects

A Cooperates

A Defects

A Cooperates	A Defects
B Cooperates	(3,3)
B Defects	(5,0)

(0,5)

(1,1)

Nash
Equilibrium



Axelrod's tournament

- Participants played the Iterated Prisoner's Dilemma (IPD)
- Tit For Tat (TFT) was the winning Strategy



Axelrod's tournament

- Participants played the Iterated Prisoner's Dilemma (IPD)
- Tit For Tat (TFT) was the winning Strategy

	A (C)	A (D)
B (C)	(3,3)	(0,5)
B (D)	(5,0)	(1,1)

TFT											
rand.											



Axelrod's tournament

- Participants played the Iterated Prisoner's Dilemma (IPD)
- Tit For Tat (TFT) was the winning Strategy

	A (C)	A (D)
B (C)	(3,3)	(0,5)
B (D)	(5,0)	(1,1)

TFT	C										
rand.	C										



Axelrod's tournament

- Participants played the Iterated Prisoner's Dilemma (IPD)
- Tit For Tat (TFT) was the winning Strategy

	A (C)	A (D)
B (C)	(3,3)	(0,5)
B (D)	(5,0)	(1,1)

TFT	C	C									
rand.	C	D									



Axelrod's tournament

- Participants played the Iterated Prisoner's Dilemma (IPD)
- Tit For Tat (TFT) was the winning Strategy

	A (C)	A (D)
B (C)	(3,3)	(0,5)
B (D)	(5,0)	(1,1)

TFT	C	C	D								
rand.	C	D	D								



Axelrod's tournament

- Participants played the Iterated Prisoner's Dilemma (IPD)
- Tit For Tat (TFT) was the winning Strategy

	A (C)	A (D)
B (C)	(3,3)	(0,5)
B (D)	(5,0)	(1,1)

TFT	C	C	D	D	D	C	C	D	C	C	
rand.	C	D	D	D	C	C	D	C	C	D	



Axelrod's tournament

- Participants played the Iterated Prisoner's Dilemma (IPD)
- Tit For Tat (TFT) was the winning Strategy

	A (C)	A (D)
B (C)	(3,3)	(0,5)
B (D)	(5,0)	(1,1)

TFT	C	C	D	D	D	C	C	D	C	C	Total = 21
rand.	C	D	D	D	C	C	D	C	C	D	Total = 26



Axelrod's tournament

- Participants played the Iterated Prisoner's Dilemma (IPD)
- Tit For Tat (TFT) was the winning Strategy

	A (C)	A (D)
B (C)	(3,3)	(0,5)
B (D)	(5,0)	(1,1)

TFT	C	C	D	D	D	C	C	D	C	C	Total = 21
rand.	C	D	D	D	C	C	D	C	C	D	Total = 26

Important:

- TFT can never get more points than the opponent
- TFT is not the best but the most robust strategy



Axelrod's tournament

- Participants played the Iterated Prisoner's Dilemma (IPD)
- Tit For Tat (TFT) was the winning Strategy

	A (C)	A (D)
B (C)	(3,3)	(0,5)
B (D)	(5,0)	(1,1)

TFT	C	C	D	D	D	C	C	D	C	C	Total = 21
rand.	C	D	D	D	C	C	D	C	C	D	Total = 26

Important:

- TFT can never get more points than the opponent
- TFT is not the best but the most robust strategy

4 Reasons TFT won:

- 1. Starts with cooperation which is nice and prevents trouble
- 2. Punishes the defection of its opponent which stimulates cooperation
- 3. Forgives the opponent when they cooperate again
- 4. TFT is a clear strategy which is easy for an opponent to anticipate



Description of Agents

The agents implemented in our model have various types of strategies



Description of Agents

The agents implemented in our model have various types of strategies:

- Always cooperating
- Always defecting



Description of Agents

The agents implemented in our model have various types of strategies:

- Always cooperating
- Always defecting
- Tit-for-tat (**TFT**): if you cooperate, I'll cooperate. Otherwise I defect



Description of Agents

The agents implemented in our model have various types of strategies:

- Always cooperating
- Always defecting
- Tit-for-tat (**TFT**): if you cooperate, I'll cooperate. Otherwise I defect
- Neural-Agents: cooperate or defect depends on the result of a neural network



Description of Agents

The agents implemented in our model have various types of strategies:

- Always cooperating
- Always defecting
- Tit-for-tat (**TFT**): if you cooperate, I'll cooperate. Otherwise I defect
- Neural-Agents: cooperate or defect depends on the result of a neural network
- String-Agents: follow a set of deterministic rules of action




Neural-Agents

Take the history(ies) of the last tournament
(0 - cooperating, 1 - defecting)

Feed the action through a forward
network

Take the output of the network
(float point number from 0 to 1) as
the tendency to defect (stochastic)



Each layer applies a linear
transform of the previous layer's
output, followed by a non-linear
activation function:

$$y_{out} = \sigma(W \cdot x + b)$$

Sigmoid function applied, to make
the output between 0 and 1



String-Agents

Histories	Decision
(CC), (CC)	C
(CC), (CD)	D
(CC), (DC)	C
(CC), (DD)	D
...	...
(DD), (DD)	D

- A lookup table that tells the agent which action to take;
- Histories: a list of (my action, opponent's action)



String-Agents


Histories	Decision
(CC), (CC)	C
(CC), (CD)	D
(CC), (DC)	C
(CC), (DD)	D
...	...
(DD), (DD)	D

- A lookup table that tells the agent which action to take;
- Histories: a list of (my action, opponent's action)
- Take all the decisions in the lookup table, we obtain a **string-representation** for the agent: CDCD...D (TFT)



Genetic Algorithm

- An algorithm to find the maxima of a system
- Inspired by the evolution in nature



Always
cooperating

Always
defecting

TFT



Genetic Algorithm

- Agents are born with different genes, which determines their behavior totally
 - E.g. String-Agents' string
 - Neural-Agents' weights and biases

TFT



Always defecting



TFT



Always cooperating






Genetic Algorithm

- Agents are born with different genes
- In each generation, the agents play against neighbors. Their **payoffs** determine their “fitness in the nature”

First round






Genetic Algorithm

- Agents are born with different genes
- In each generation, the agents play against neighbors. Their **payoffs** determine their “fitness in the nature”

Subsequent rounds






Genetic Algorithm

- Agents are born with different genes
- In each generation, the agents play against neighbors. Their **payoffs** determine their “fitness in the nature”

After last round of the generation





Genetic Algorithm

- Agents are born with different genes
- In each generation, the agents play against neighbors. Their **payoffs** determine their “fitness in the nature”
- The fitter agents will survive, whereas the weaker ones die

After last round of the generation



Dominant
Species


Disadvantageous
Species



Genetic Algorithm

- The fitter agents will survive, whereas the weaker ones die
- The fittest agents reproduce with each other

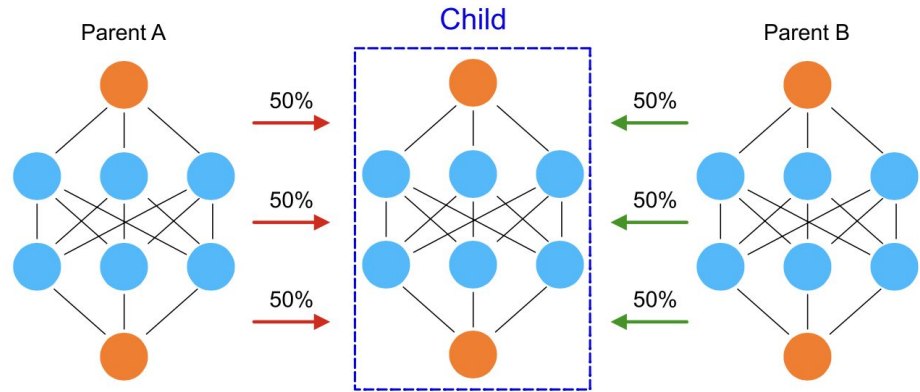
After last round of the generation





Genetic Algorithm


- The fitter agents will survive, whereas the weaker ones die
- The fittest agents reproduce with each other
 - Parents cross over their genes (or swap weight matrices) to create the children





Genetic Algorithm


- The fittest agents reproduce with each other
 - Parents cross over their genes (or swap weight matrices) to create the children
 - Random mutations can happen to children's genes





Genetic Algorithm


- The fittest agents reproduce with each other
 - Parents cross over their genes (or swap weight matrices) to create the children
 - Random mutations can happen to children's genes
 - The child replace the weakest neighbor






Genetic Algorithm

- The fittest agents reproduce with each other
 - Parents cross over their genes (or swap weight matrices) to create the children
 - Random mutations can happen to children's genes
 - The child replace the weakest neighbor
- Redo the former process for next generations





Genetic Algorithm





Project goals

Can the Agents learn Tit for Tat


Investigating the behavior of Agents in different settings

Analyzing the formation of cooperation





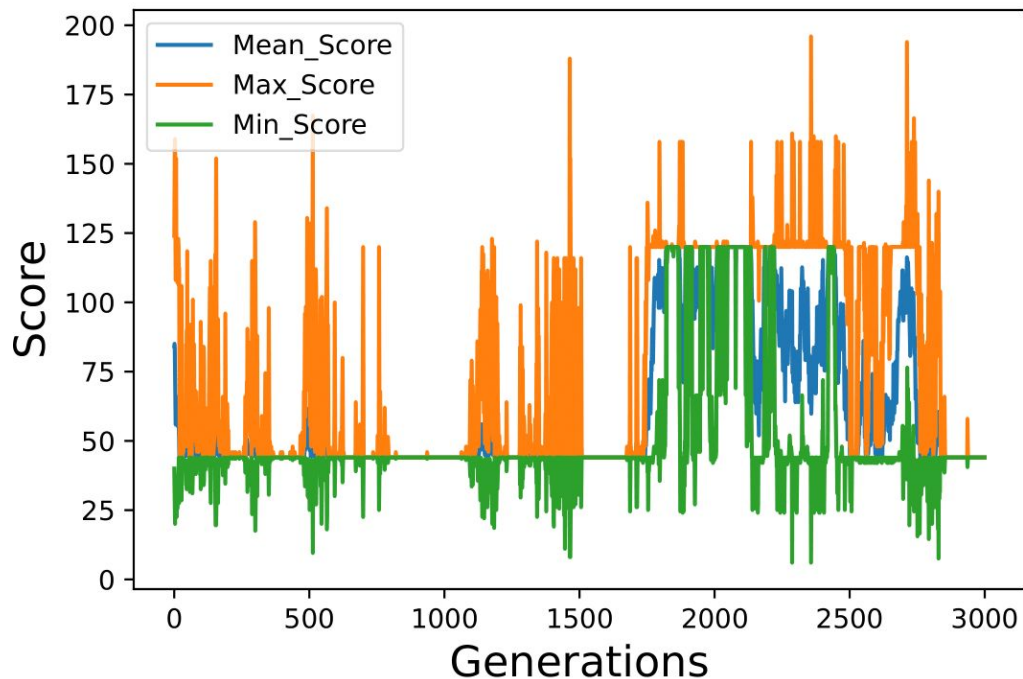
Ring Structure





Ring Structure


- 20 Agents
- Nearest-Neighbor-Interaction
- 20 Rounds of PD
- Global-Genetic-Algorithm





Ring Structure - Feature Vectors


- **FV0:** Output, given C of opponent
 - FV0 = 0: Cooperating
 - FV0 = 1: Defecting
- **FV1:** Output, given D of opponent
 - FV1 = 0: Forgiving
 - FV1 = 1: Punishing
- Phase of cooperation from 1700 to 2500





Ring Structure - Feature Vectors


- Starts with Tit for tat
- Always Cooperating
- Defectors spread
- Tit for Tat again







Ring - Influence from outside

What we already know
For Neural-Agents?



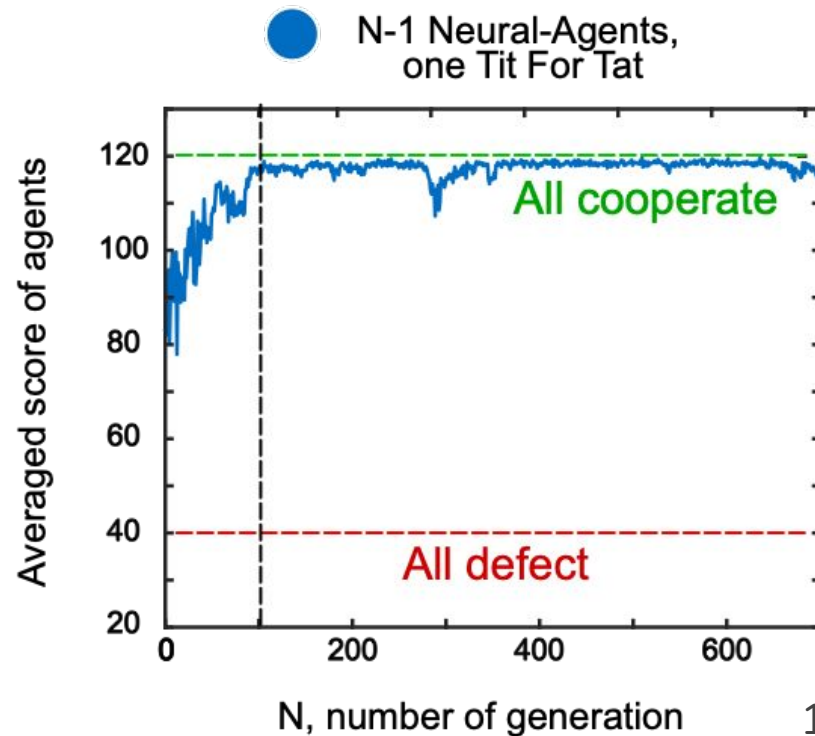

Ring - Influence from outside

What we already know
For Neural-Agents?




Ring - Influence from outside


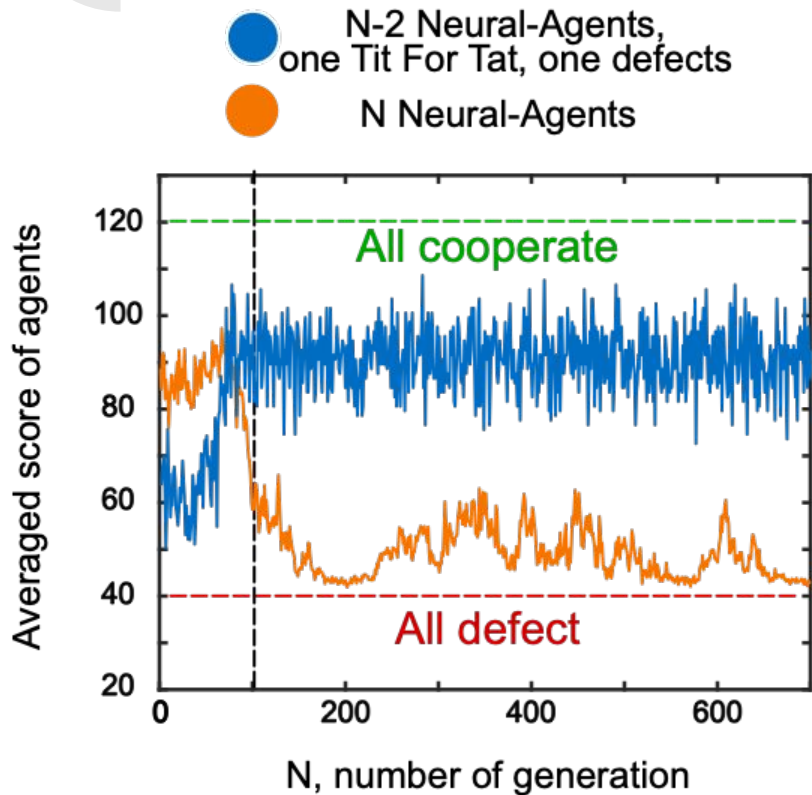
What we already know
For Neural-Agents?



Ring - Influence from outside

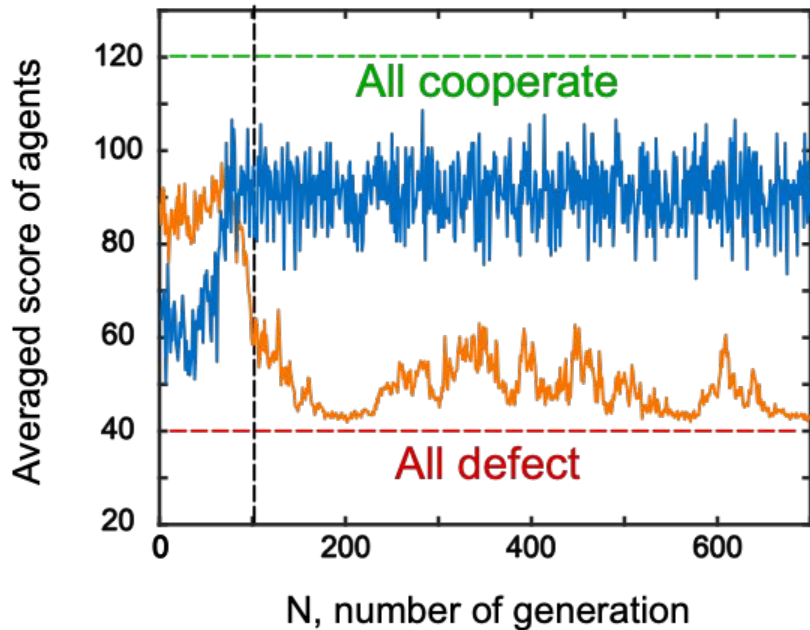


Ring - Influence from outside




Ring - Influence from outside

- N-2 Neural-Agents,
one Tit For Tat, one defects
- N Neural-Agents

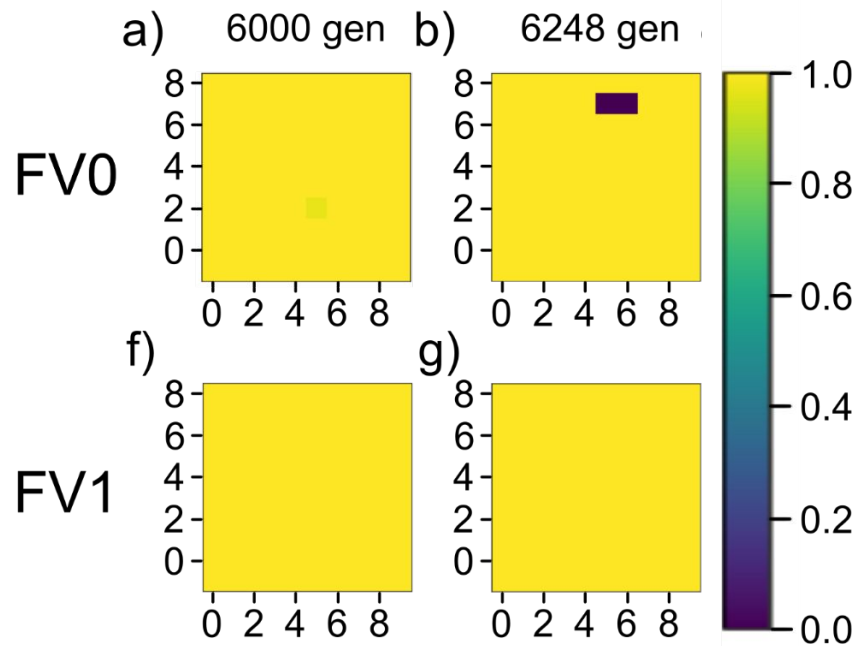


! Let's go to 2D !



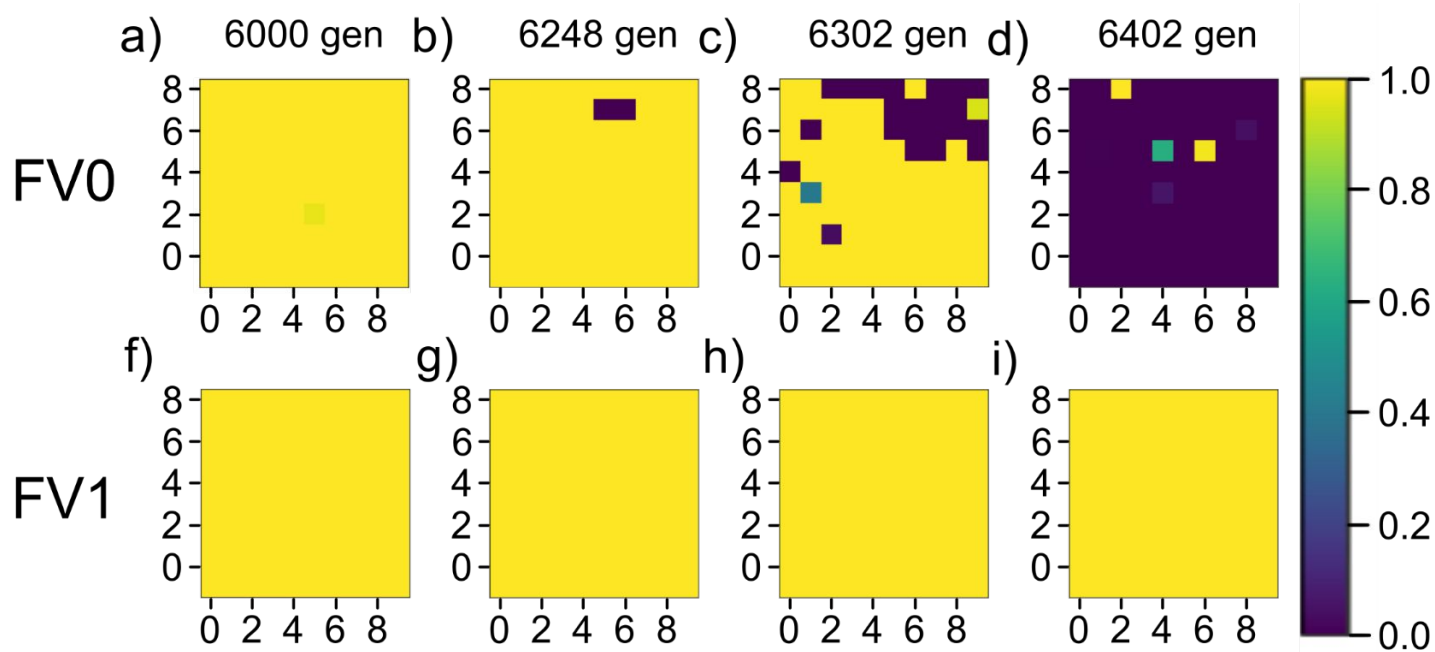


Grid - Formation of Cooperation



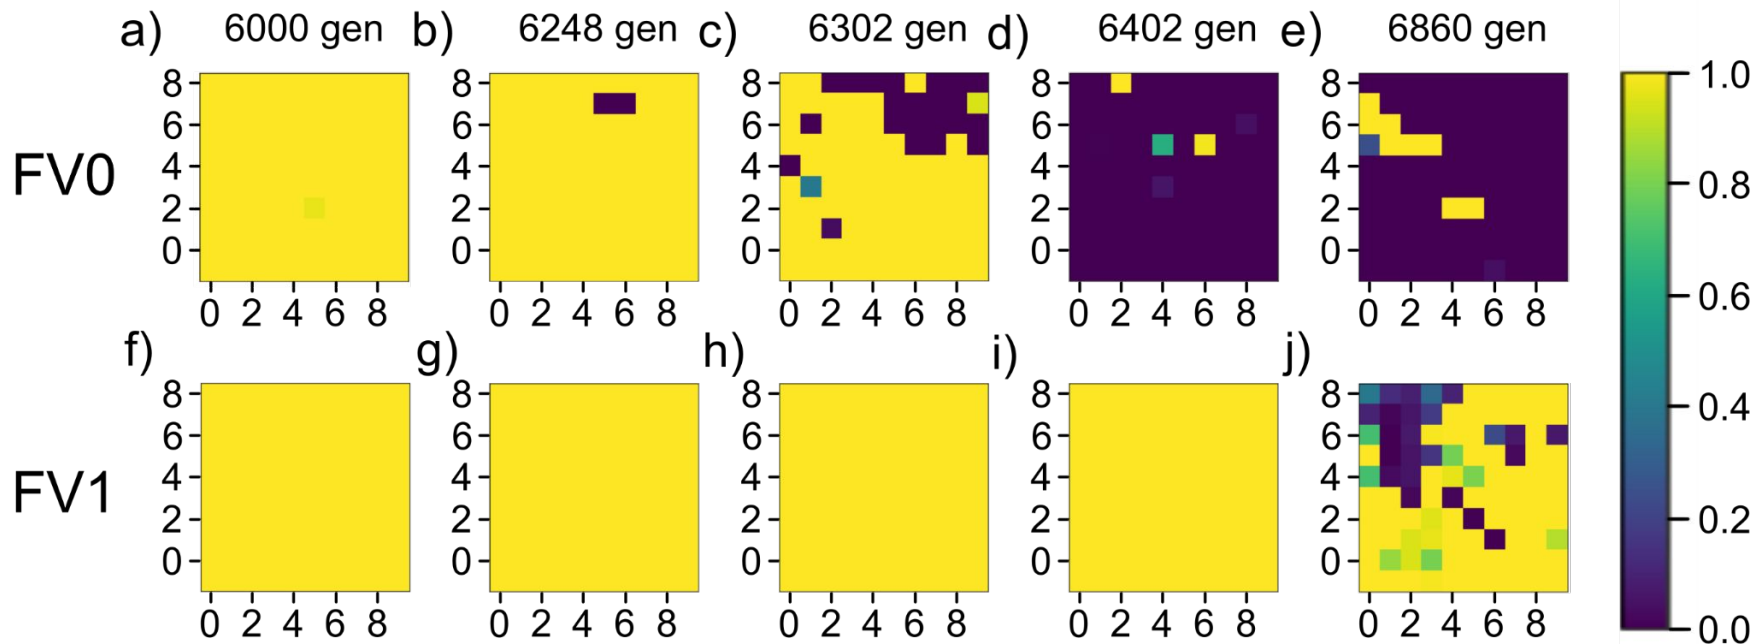


Grid - Formation of Cooperation





Grid - Formation of Cooperation





Config.py

```
# 0 cooperating, 1 defecting
# Map (my_action, other_action) to payoff
PAYOFF_MAP = [
    [3, 0],
    [5, 1],
]

# Neighbors of each agent: can be 8 (everyone around), 4 (left right top bottom), or 2 (left right)
NEIGHBOR_TYPE = 4
NEIGHBOR_RADIUS = 1 # interaction radius

NUM_SUBSTEPS = 10 # rounds in a generation

# Genetic Algorithm parameters
USE_LOCAL_GA = True # whether to use the localized version of GA
FITNESS_MULTIPLIER = 2 # parameter used for scaling fitness function, see P15 of the java manual
EPS = 1e-3 # for fitness scaling
MUT_PROB = 0.2 # mutation probability for NeuralAgent and StringAgent
MUT_STRENGTH = 1.0 # how strong to perturb the weights of NeuralAgent
TFT_REPRODUCABLE = False # does TFT participate in the reproduction
NEURAL_REPRODUCABLE = True # does NeuralAgent participate in the reproduction
STRING_REPRODUCABLE = True # does StringAgent participate in the reproduction

# Network topology for NeuralAgent
# e.g. [6, 1] means a network with 6 inputs (i.e. memory size 3), fully connected to the output
DEFAULT_NEURAL_STRUCTURE = [6, 1]

# StringAgent's memory size
MEM_LEN = 2
```

```
# StringAgent's memory size
MEM_LEN = 2

# Settings for PDModel
DEFAULT_WIDTH = 10
DEFAULT_HEIGHT = 10
TORUS_GRID = True
CANVAS_DX = 30

# Random seeds
MESA_SEED = 4
NUMPY_SEED = 3

# The way agents are rendered in the We
# 'agent_type': the colors indicate whi
# 'defecting_ratio': visualize how defe
# 'inherited_attr': visualize who are t
VISUALIZE_GRID_TYPE = 'inherited_attr'
```



Website application via Mesa package

Evolution of Prisoner's Dilemma About

Number of steps in each generation



Fitness function type

score ▾

Agent type

mixed ▾


Neighbor type

4 ▾

Frames Per Second



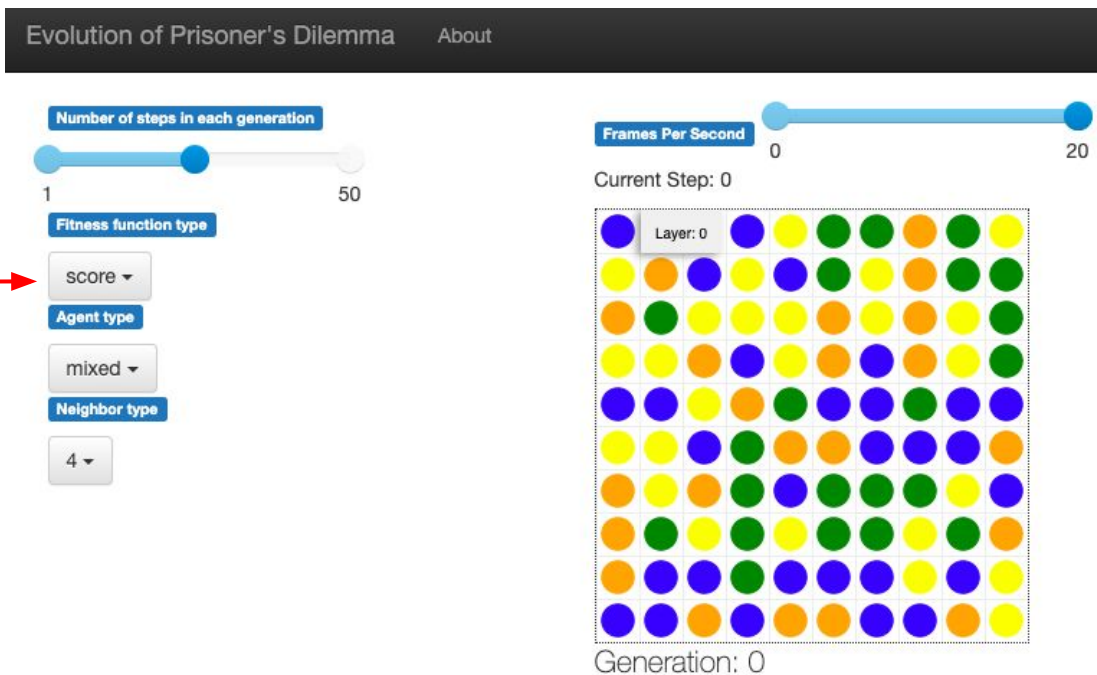
Current Step: 0



Generation: 0



Website application via Mesa package



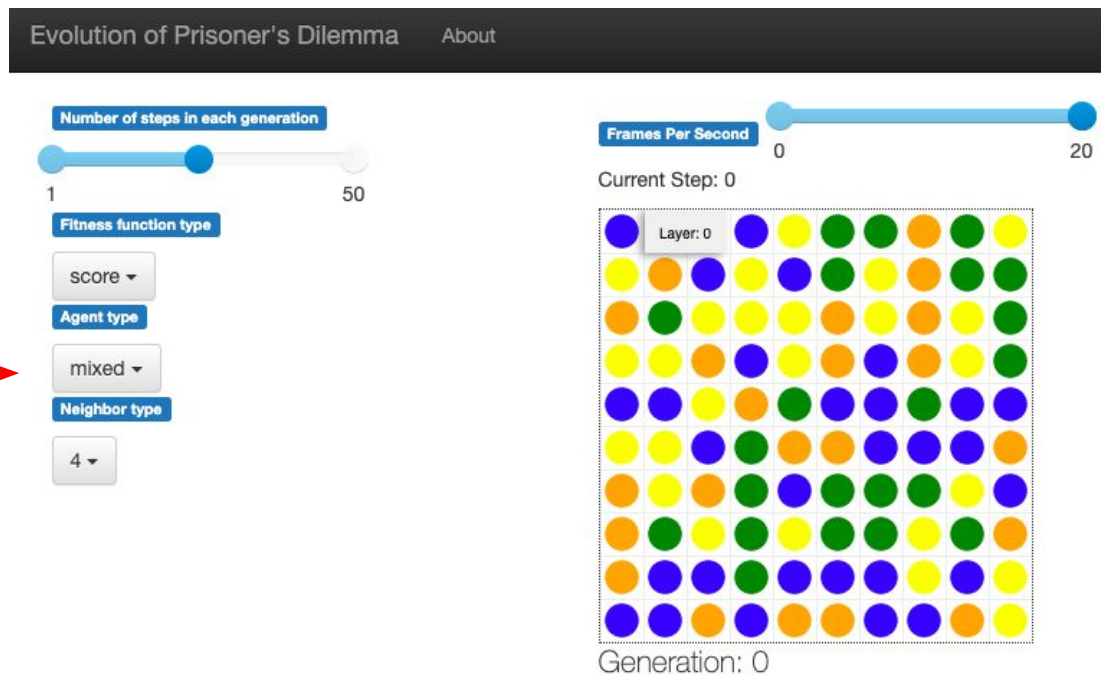


Website application via Mesa package

Agent type

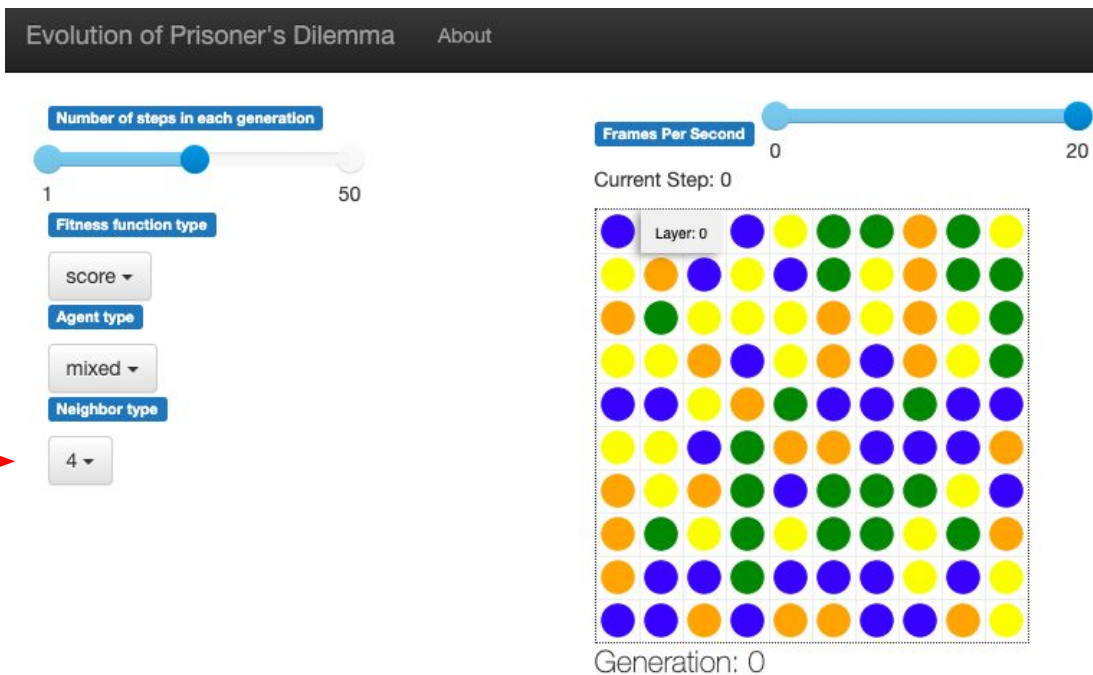


neural
string
tit_for_tat
simple
mixed





Website application via Mesa package



Number of
neighbors



8


4

2



Website application via Mesa package

Generation: 109 Fittest feature vector: 0.4918, 0.6789



Distribution of
agents



Execution panel

Start Step Reset

Distribution of
scores





Results

Agents have to learn TFT for cooperation

Cooperation could be achieved by adding punishers

Local interaction + reproduction are important for cooperation

Larger Agent memory leads to stable cooperation



Questions






References

1. R. M. Axelrod and W. D. Hamilton, *The Evolution of Cooperation*. Basic Books, 1984, isbn: 0-465-02121-2.
2. R. Axelrod, *The Evolution of Strategies in the Iterated Prisoner's Dilemma*. Lawrence Davis, London: Pitman, and Los Altos, CA: Morgan Kaufman, 1987, pages 32–41.
3. A. Errity, *Evolving Strategies for the Prisoner's Dilemma*, 2003.
4. J. H. Holland, *Genetic Algorithms*, 1. Scientific American, a division of Nature America, Inc., 1992, volume 267, pages 66–73. url: <http://www.jstor.org/stable/24939139>.


Memory-length tests

● Max ● Mean ● Min

a) Scores for length-1 memory




b) Mean feature vectors for length-1 memory




Memory-length tests

● Max ● Mean ● Min

c) Scores for length-2 memory




d) Mean feature vectors for length-2 memory



Memory-length tests


● Max ● Mean ● Min

e) Scores for length-3 memory




f)

Mean feature vectors for length-3 memory



Connection to society?

- Public Good Game



- Examples: Tax evasion, climate change, cleaning shared apartment, Cross-Code-Checking
- Defectors lead to tragedy of the common
- Punishment and Reputation key for avoiding the tragedy of the commons